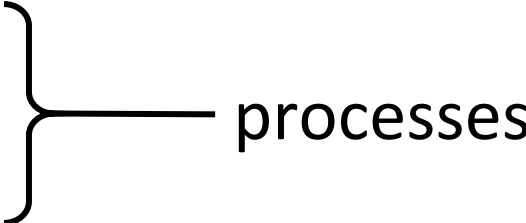


Processes

Process concept

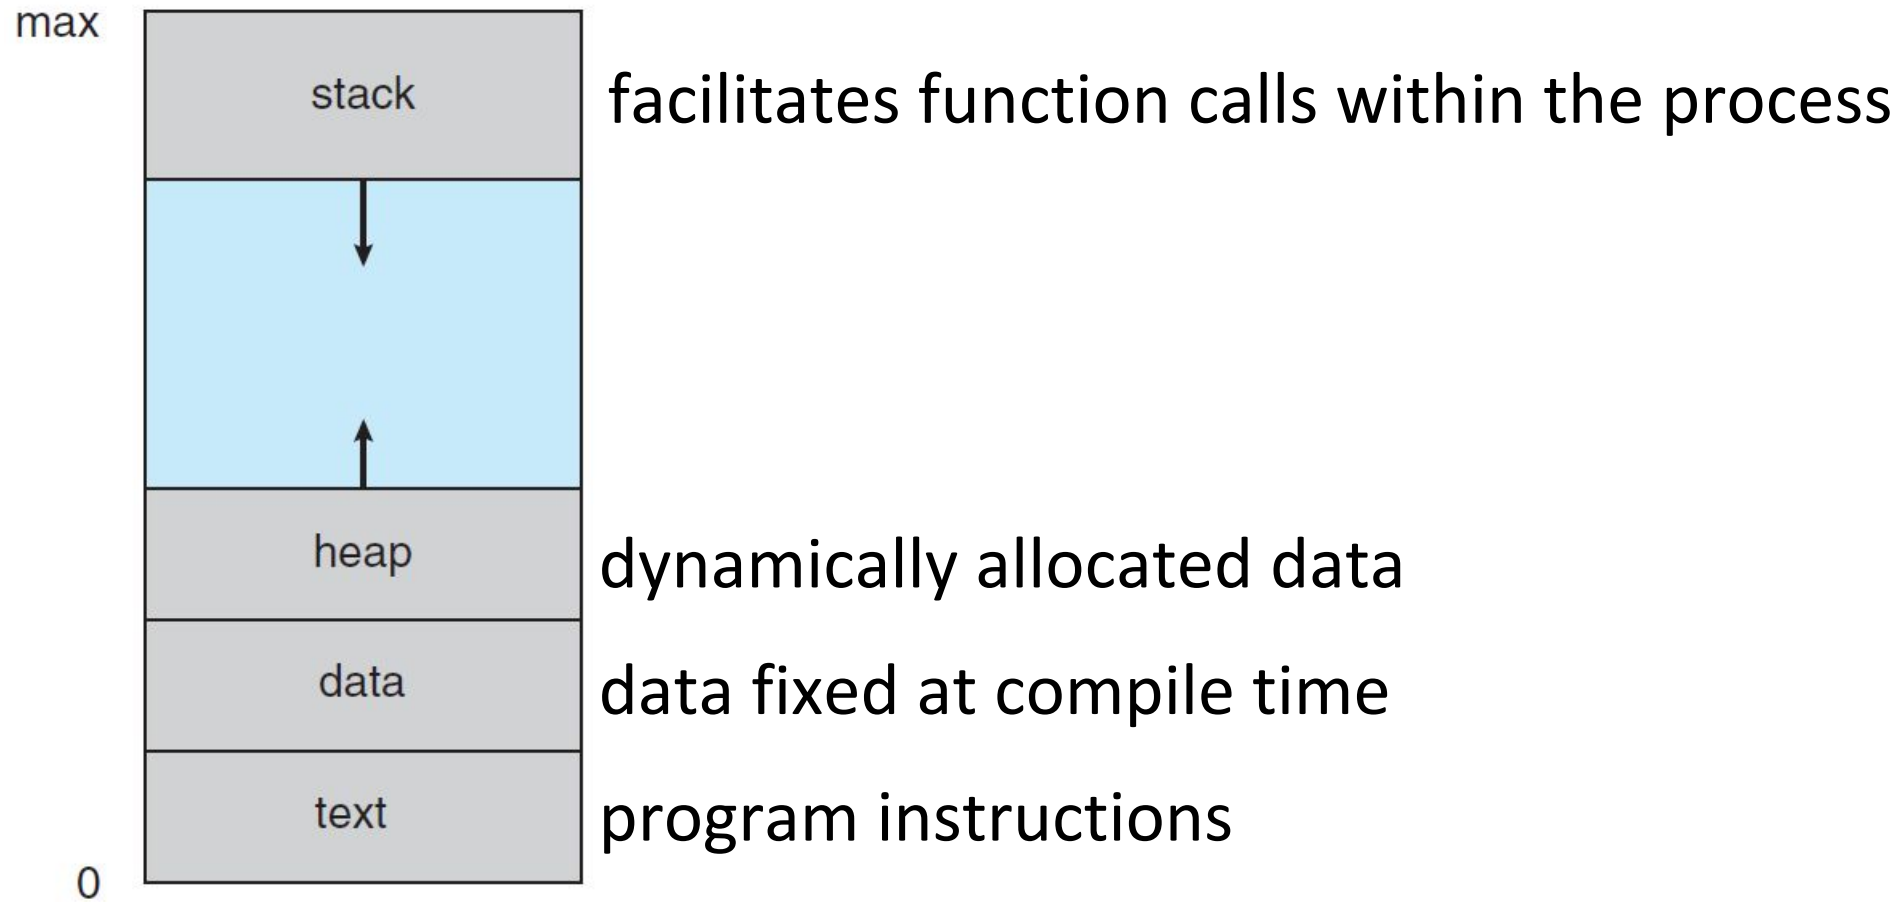
- A process is a program in execution
 - Program is a sequence of instruction in a file (executable file) on the disk
 - Program is passive
 - Process is “alive”
 - A single program can have multiple process
 - A single process can spawn multiple processes
 - Batch systems – jobs
 - Time-sharing systems – user programs or tasks
- 
- processes

Process concept

- A process includes

- program counter – address of next instruction
 - contents of CPU registers
 - text section
 - stack
 - heap
 - data section
-
- The diagram uses curly braces to group the components of a process. A brace on the right side groups the 'program counter – address of next instruction' and 'contents of CPU registers' under the label 'CPU'. Another brace on the left side groups the 'text section', 'stack', 'heap', and 'data section' under the label 'Memory'.
- CPU
- Memory

Process concept – memory



Single program with multiple process

- Each process has distinct
 - program counter
 - CPU register contents
 - stack
 - heap
 - data section
- The processes share the text section only

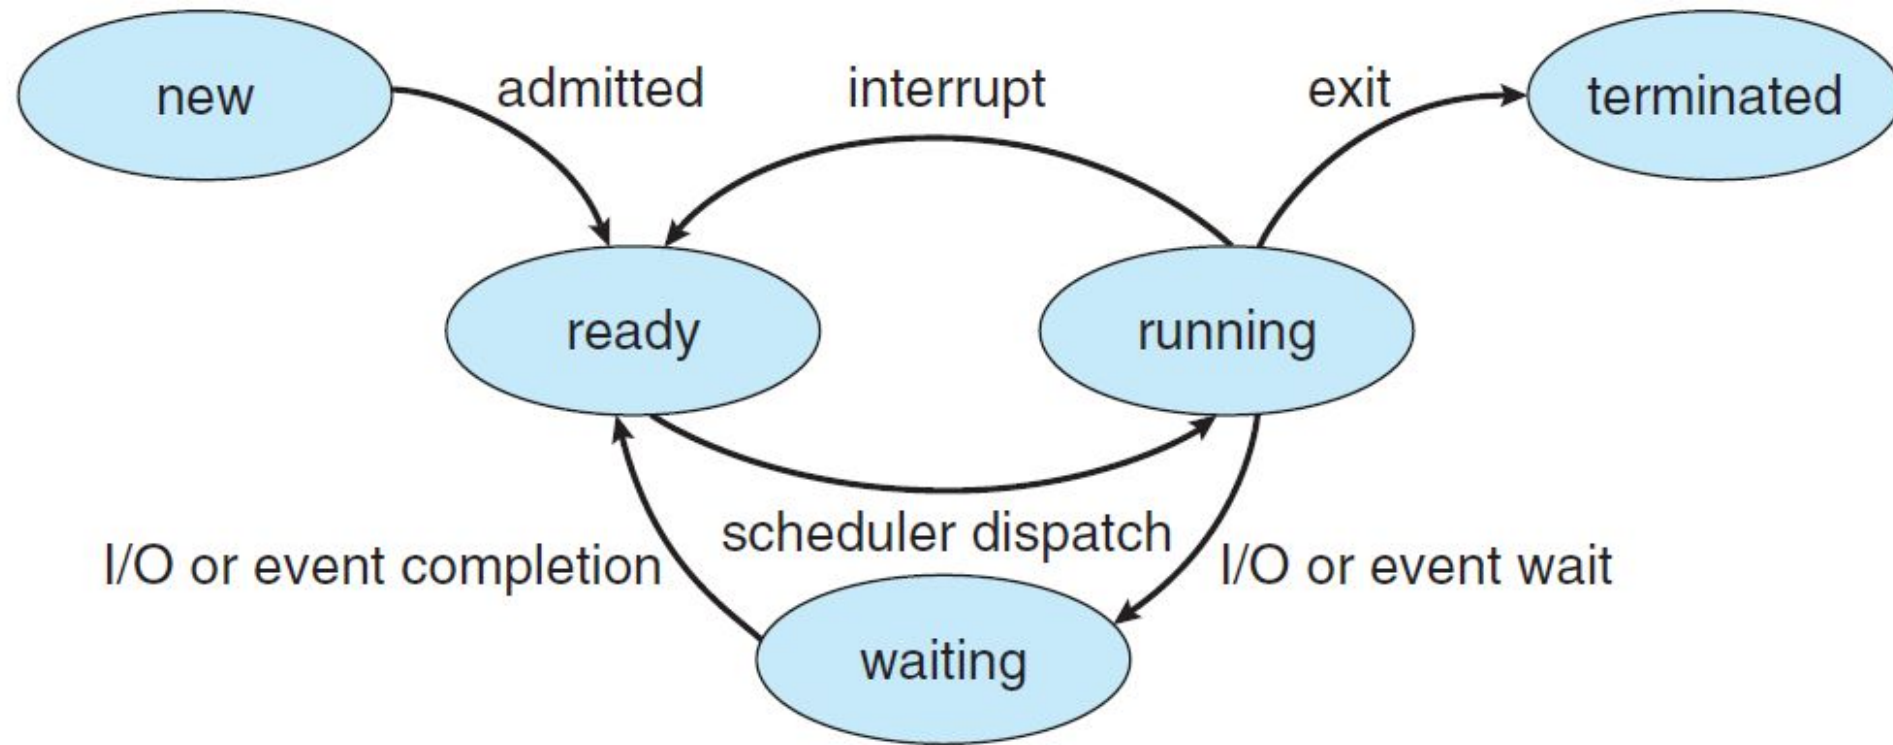
Process as an execution environment

- A process can act as an execution environment for another process
 - Emulation
- Example: JVM
 - java myProgam**
 - **java** runs the JVM process which is responsible interpreting **myProg** (a processe also in memory) instructions into native instructions on the CPU

Process states

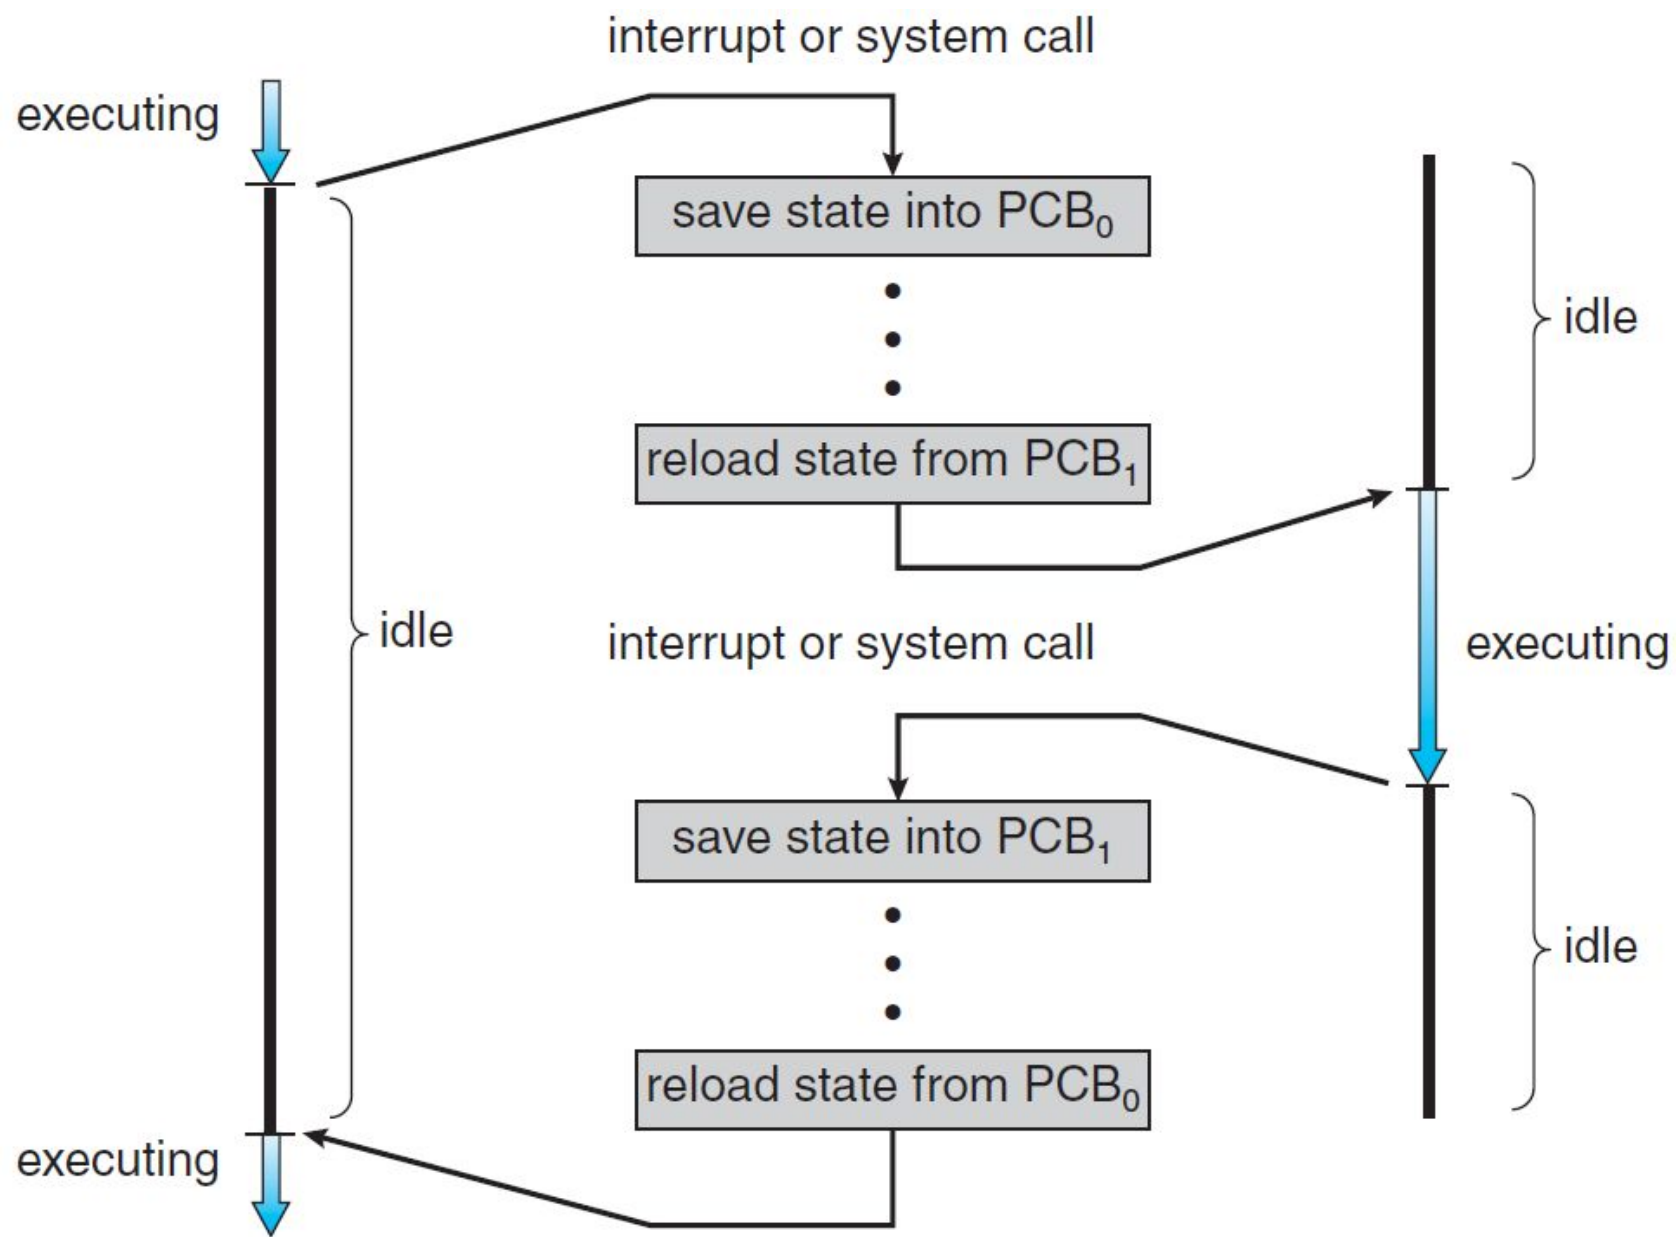
- **new**
- **running** – currently executing in the CPU
- **waiting** – waiting for some event
- **ready** – waiting for a turn to use the CPU
- **terminated**

Process states



Process control block (PCB)

- Also called **task control block**
- OS maintains a PCB for each process
- PCB contains information about the process including
 - process state
 - CPU register contents
 - CPU scheduling information
 - Memory management info – e.g. values of base address and limit registers etc
 - Accounting info: real time used, time limits, process numbers, etc
 - I/O status information: lists of allocated devices, open files, etc



Threads

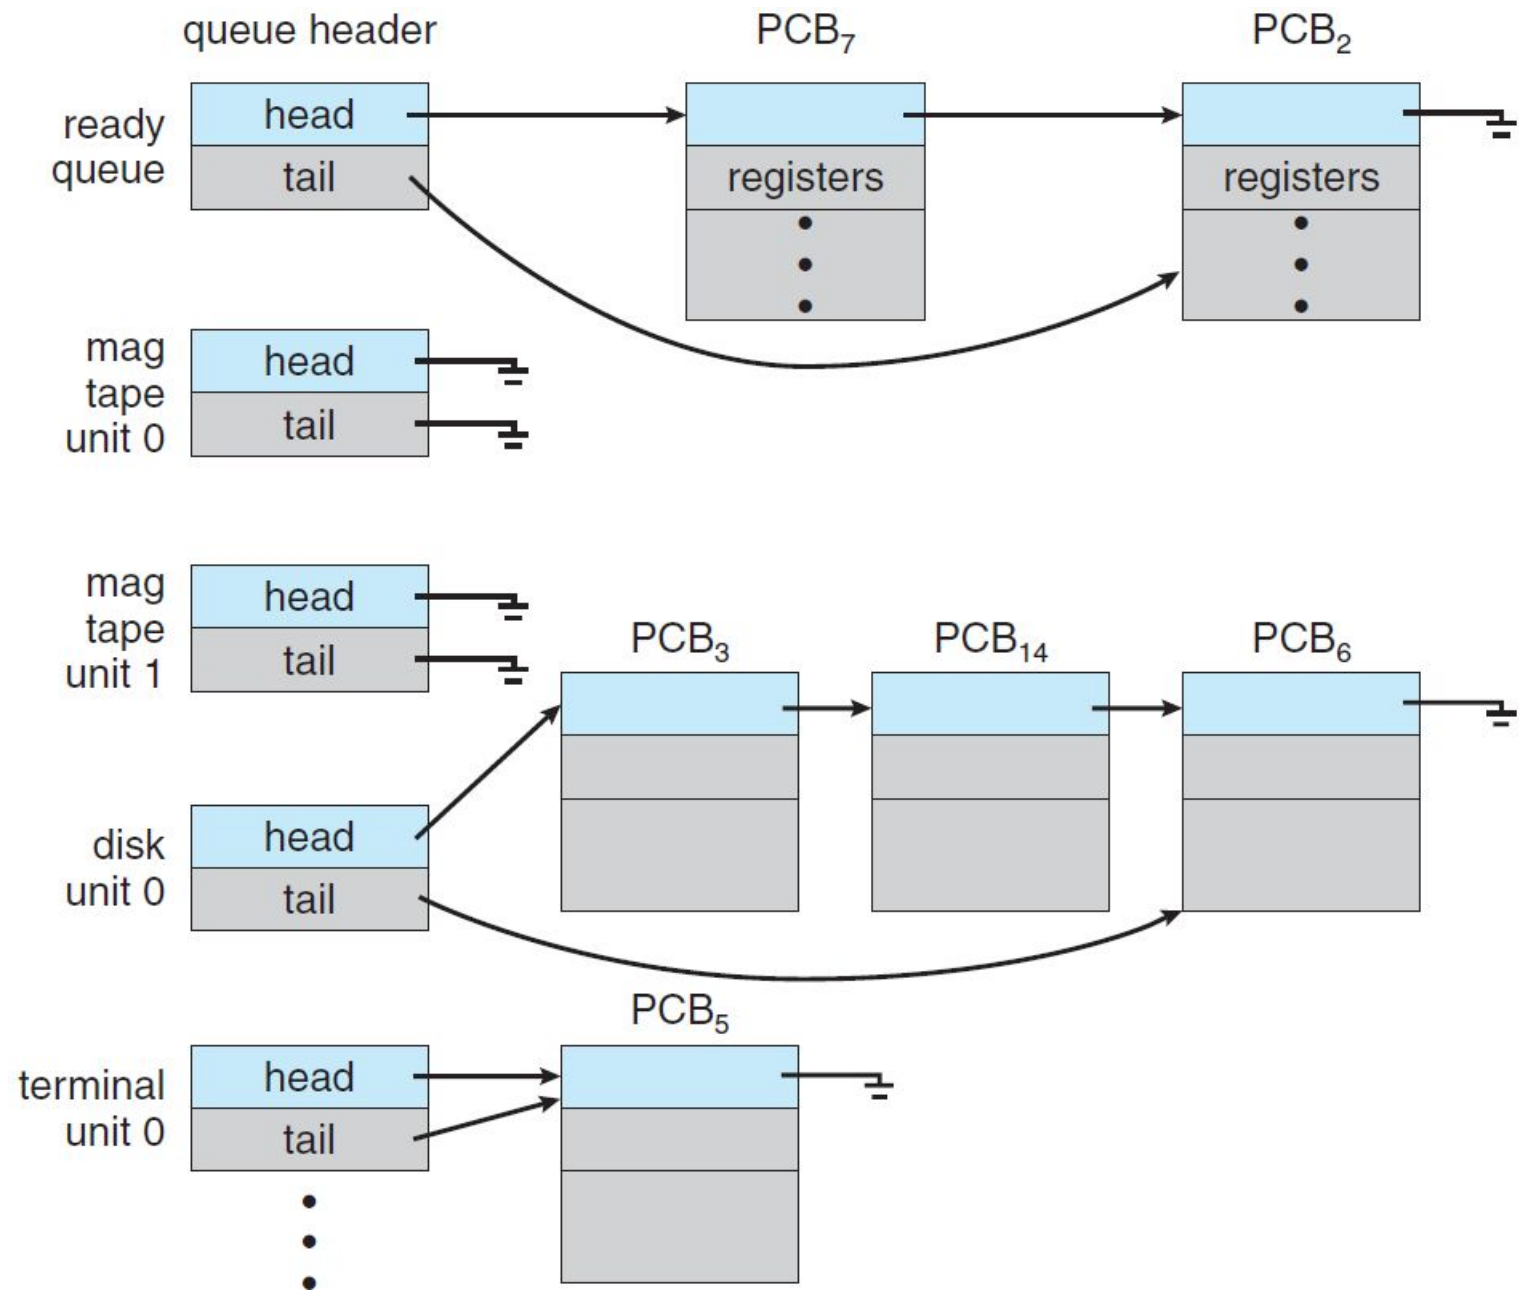
- A process is a single instance of a program in execution
- A process can be multithreaded – multiple separate sequences of instructions, each requiring a separate program counter, etc
- PCBs for multithreaded process are expanded to include additional information
- Multithreaded processes can take advantage of multiprocessor system to run threads in parallel

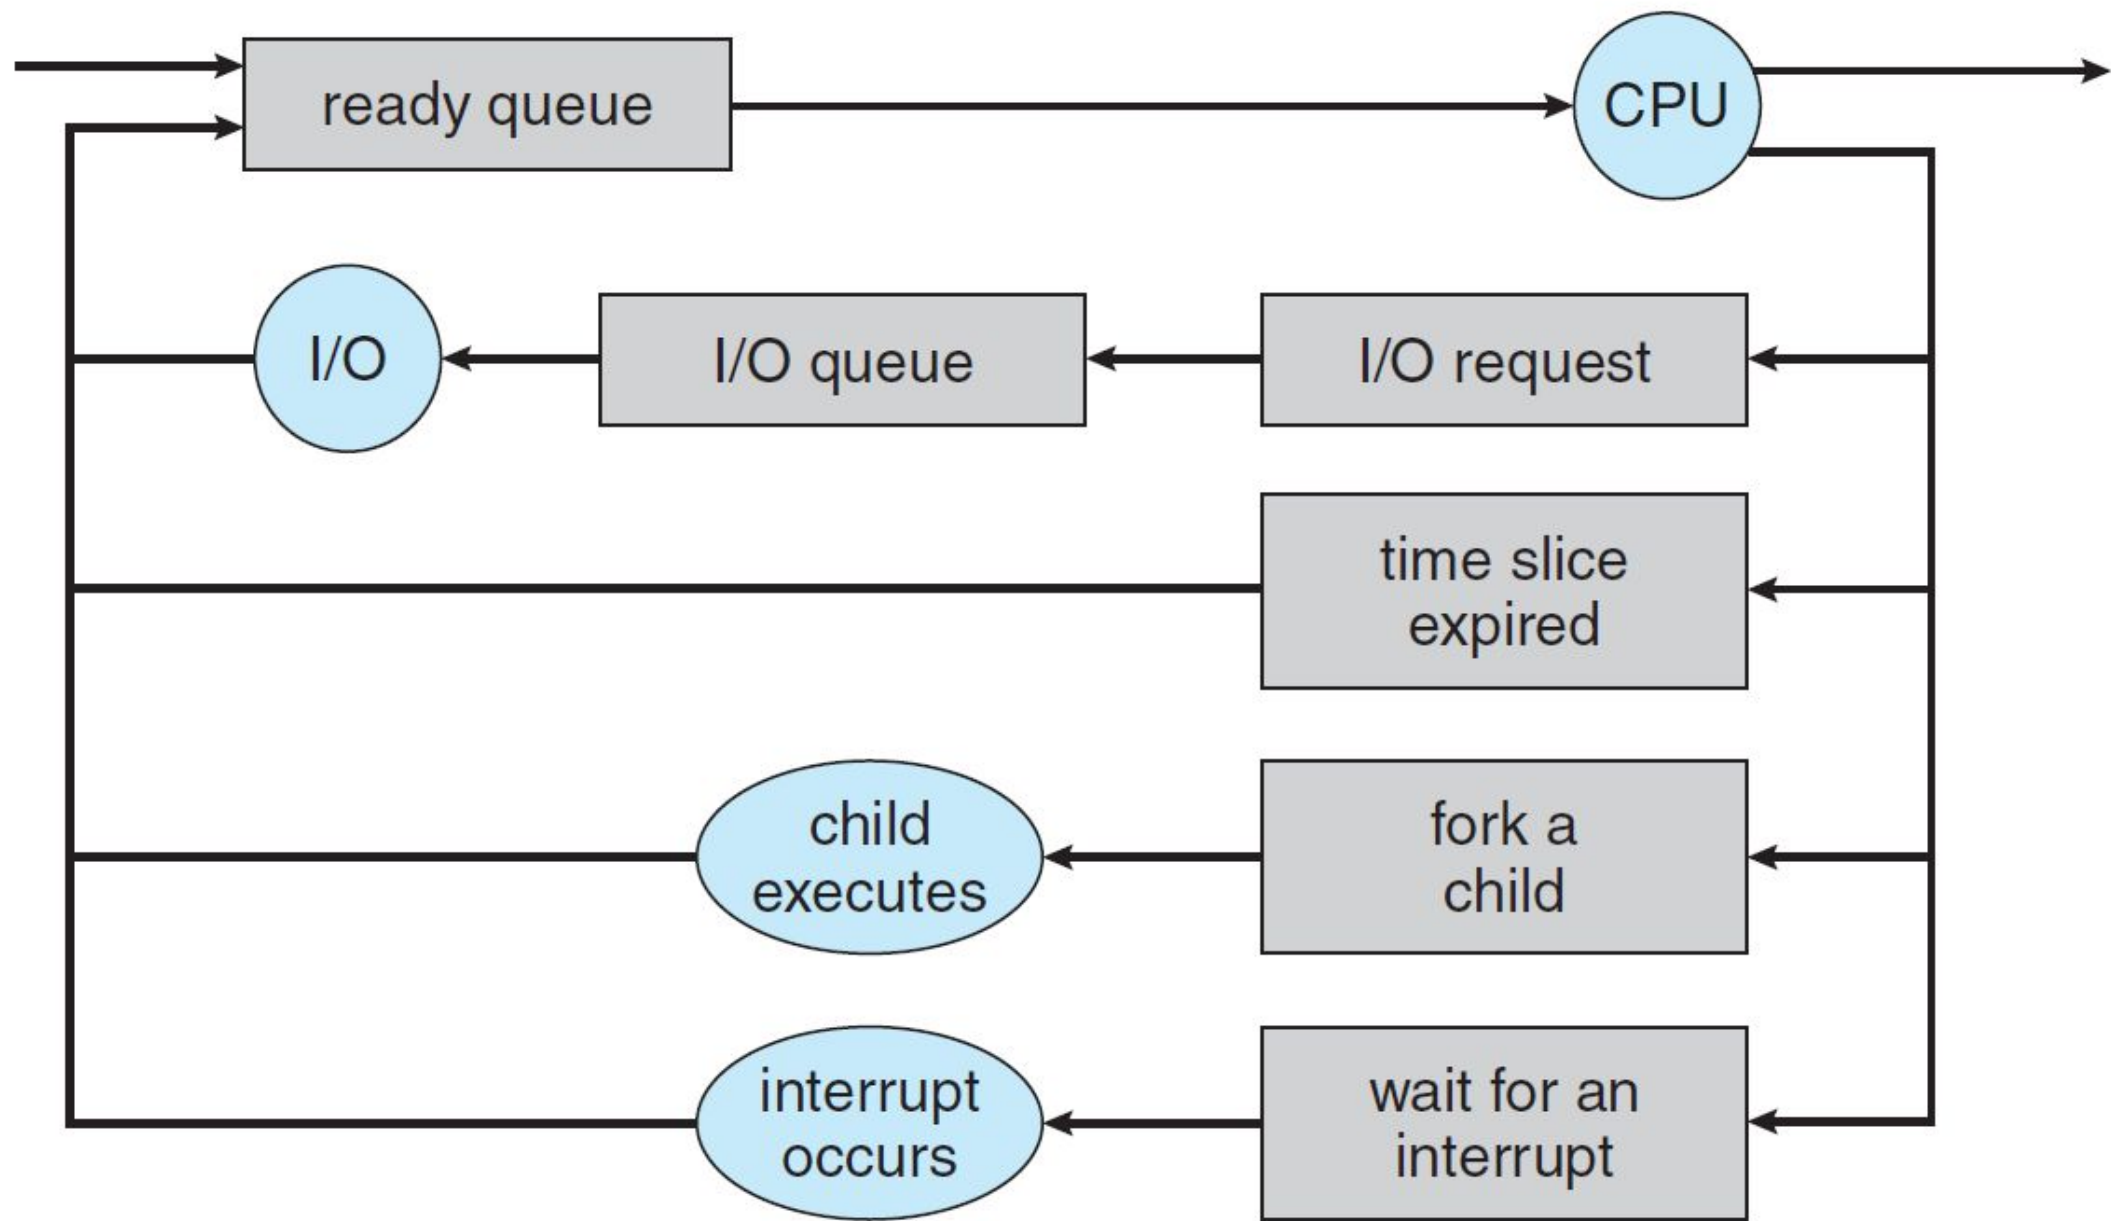
Process scheduling

- A must in a multiprogramming environment
- Multiprogramming promotes high resource utilization
 - E.g. Let another process run while the current one is waiting for IO
- Allows user to run multiple applications “simultaneously”
- Allows multiple users to use a system “concurrently”
- A single CPU can only execute one process at any given time
 - The other processes must wait until **process scheduler** select one for execution

Scheduling queues

- OS maintains several queues for processes waiting for some resource
 - Job queue – newly created processes
 - Ready queue
 - Device queues
- Each queue is implemented as a linked-list of PCBs
 - Queue header has pointers to first and last PCB





Schedulers

- Long-term scheduler (job scheduler)
 - For jobs spooled on the disk
 - Selects jobs from disk for loading into memory
 - Controls **degree of multiprogramming** – number of processes in memory
 - Low frequency of use
- Short-term scheduler (CPU scheduler)
 - Selects, from the ready queue, a process to run on the CPU
 - Called frequently – typically, every few milliseconds
 - Must decide very quickly, otherwise it will waste CPU cycles

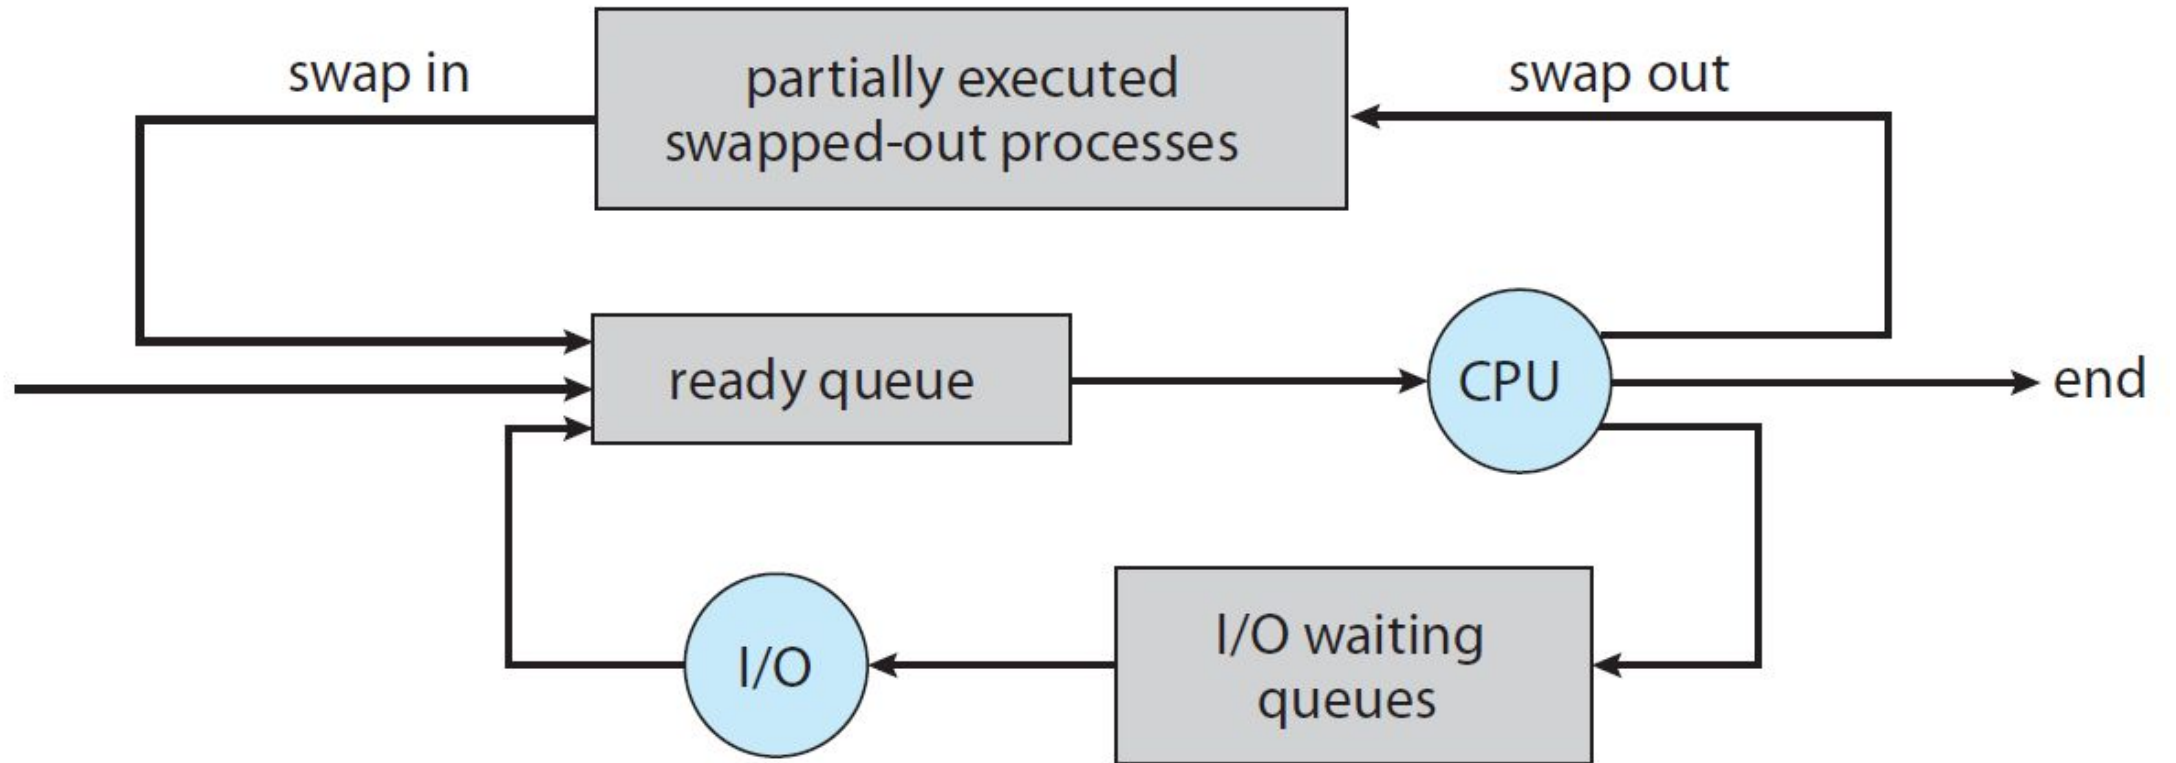
Long-term scheduler

- I/O bound processes – processes that spend most of their time doing I/O
- CPU bound processes – processes that are computation intensive
- long-term scheduler must balance between the two
 - if too many CPU bound processes – I/O devices will be idle
 - if too many I/O bound processes – CPU will be mostly idle
- Some OSs (e.g. Windows and Unix) don't have long-term scheduler
 - new processes immediately join the ready queue
 - degree of multiprogramming is controlled by self adjusting nature of users

Medium-term scheduler

- In OSs that support swapping
- Invoked when the degree of multiprogramming becomes too high for the available amount of memory
 - Some processes are swapped out and later swapped back in when memory becomes more available
- Might also be used to balance the process mix

Medium-term scheduler in context



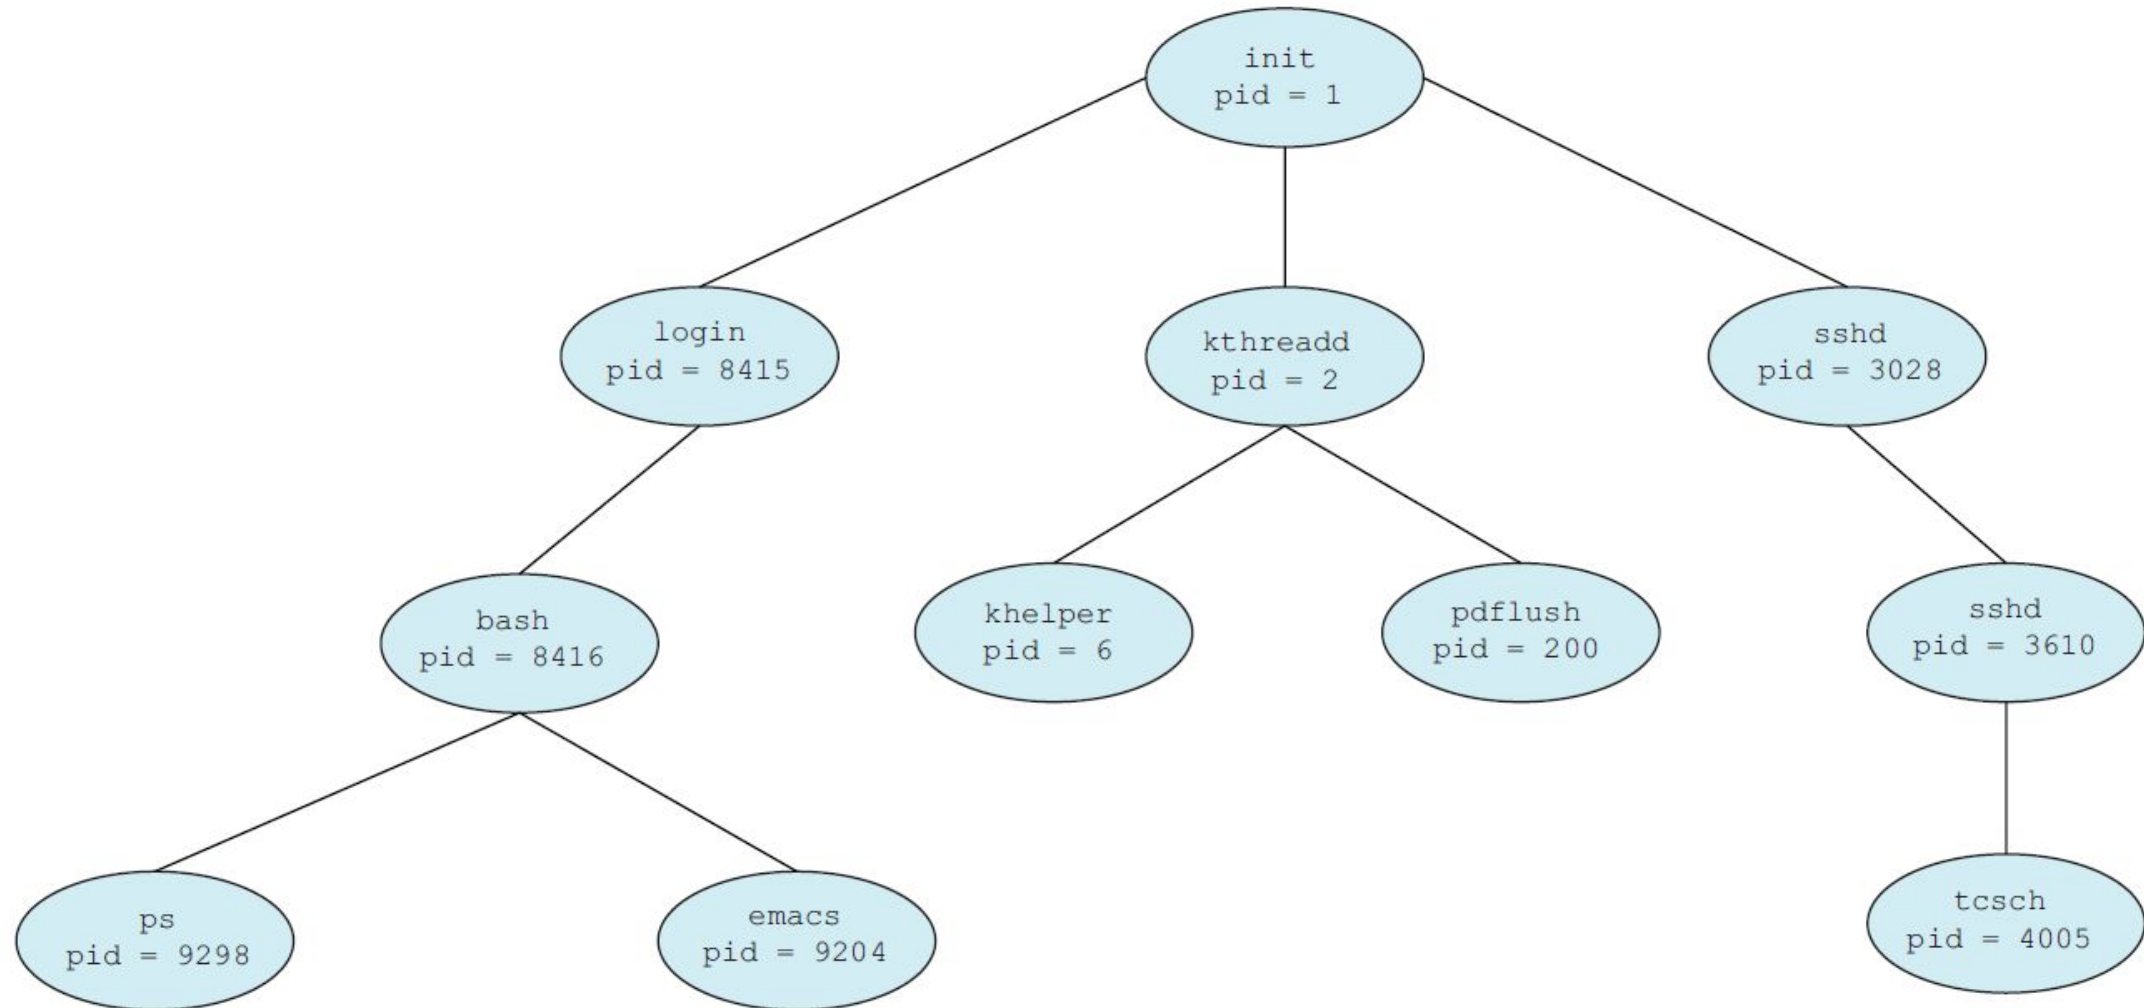
Context switch

- Context = the contents of a PCB
- context switch happens at each interrupt or system call
 - state save of the outgoing process
 - state restore of the incoming process
- Speed of a context switch depend on
 - complexity of the OS
 - the hardware architecture
 - E.g. some CPU have register sets – can simply switch to another register set instead of copying register values to memory (in PCB)

Operations on Processes - Creation

- Parent/child processes
 - Process tree
- Each process is identified by a process ID (PID)
- After a process creates a new process, either
 - Concurrent execution of parent process and its children
 - Parent process waits for the child process to complete execution
- Address space possibilities
 - Child process duplicates parent's address space
 - Child process loads a new program

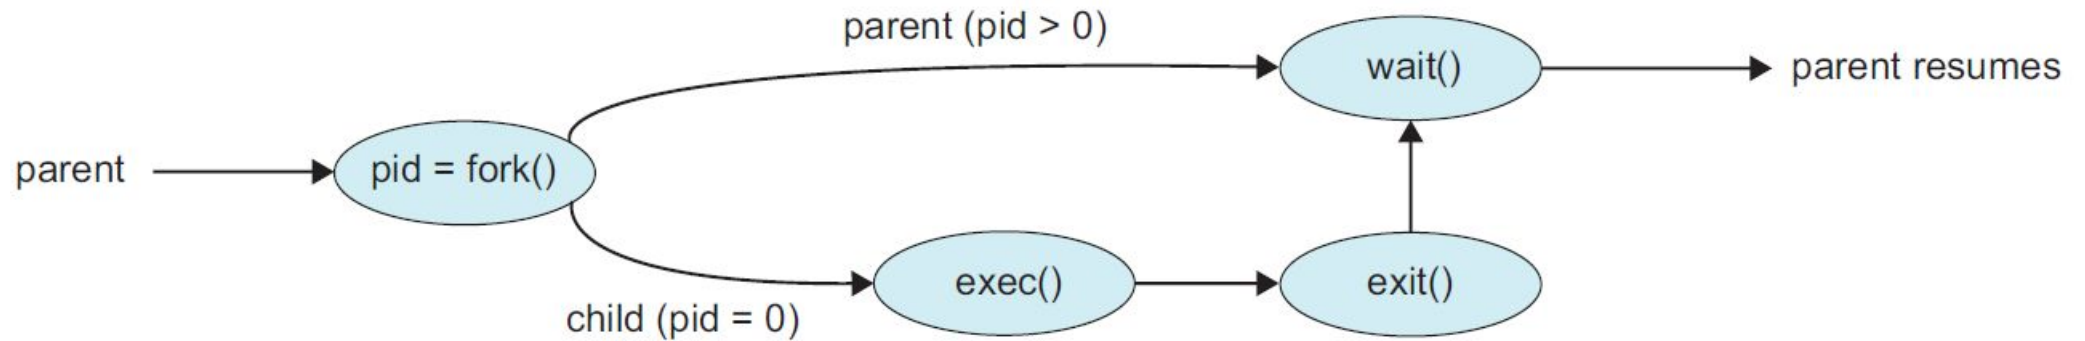
Process tree example: UNIX



Process creation example: UNIX

- **fork()** – system call to create a child process
 - Creates a child process with a copy of the parent's address space
 - parent and child run concurrently, unless parent chooses to call **wait()**
 - **fork()** in parent returns PID of the newly created process
 - **fork()** in child return 0
- **exec()** – loads a new program in a process

Process creation example: UNIX




```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>
int main(){
    pid_t pid;

    /* fork a child process */
    pid = fork();

    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete");
    }

    return 0;
}
```

Process creation example: Windows

- **CreateProcess()** – system call to create process
 - has 8 parameters
 - includes the program that must be run in the process

```

#include <stdio.h>
#include <windows.h>
int main(VOID){
    STARTUPINFO si; PROCESS_INFORMATION pi;
    /* allocate memory */
    ZeroMemory(&si, sizeof(si)); si.cb = sizeof(si); ZeroMemory(&pi, sizeof(pi));
    /* create child process */
    if (!CreateProcess(NULL, /* use command line */
        "C:\\WINDOWS\\system32\\mspaint.exe", /* command */
        NULL, /* don't inherit process handle */
        NULL, /* don't inherit thread handle */
        FALSE, /* disable handle inheritance */
        0, /* no creation flags */
        NULL, /* use parent's environment block */
        NULL, /* use parent's existing directory */
        &si, &pi)){
        fprintf(stderr, "Create Process Failed"); return -1;
    }
    WaitForSingleObject(pi.hProcess, INFINITE); /* parent to wait for child */
    printf("Child Complete");
    /* close handles */ CloseHandle(pi.hProcess); CloseHandle(pi.hThread);
}

```

Operations on process: termination

- `exit()` – system call made by process to request OS to delete it
 - May return a value (typically an integer) representing status to parent process – assuming parent called `wait()` on the process
 - Deallocate all resources
 - Can be called explicitly
 - Can also be called implicitly – return statement
- `TerminateProcess()` – system call to terminate a given process
 - Windows OS
 - A process can only terminate its children

Operations on process: termination

- Use cases for terminating a process
 - Child exceed usage of an allocated resource
 - Task assigned to child is no longer needed
 - Parent process is exiting (**Cascading termination**) - some OSs do not allow orphans
- Zombie processes – a process that has terminated, but, its parent has not yet called wait
 - Its entry in the process table still exists – so that the parent is able to get the return status
 - Can only be deleted from the process table if a parent calls wait()

Operations on process: termination

```
/* exit with status 1 */  
exit(1);
```

Child process

```
pid_t pid;  
int status;  
pid = wait(&status);
```

Parent process

Operations on process: termination

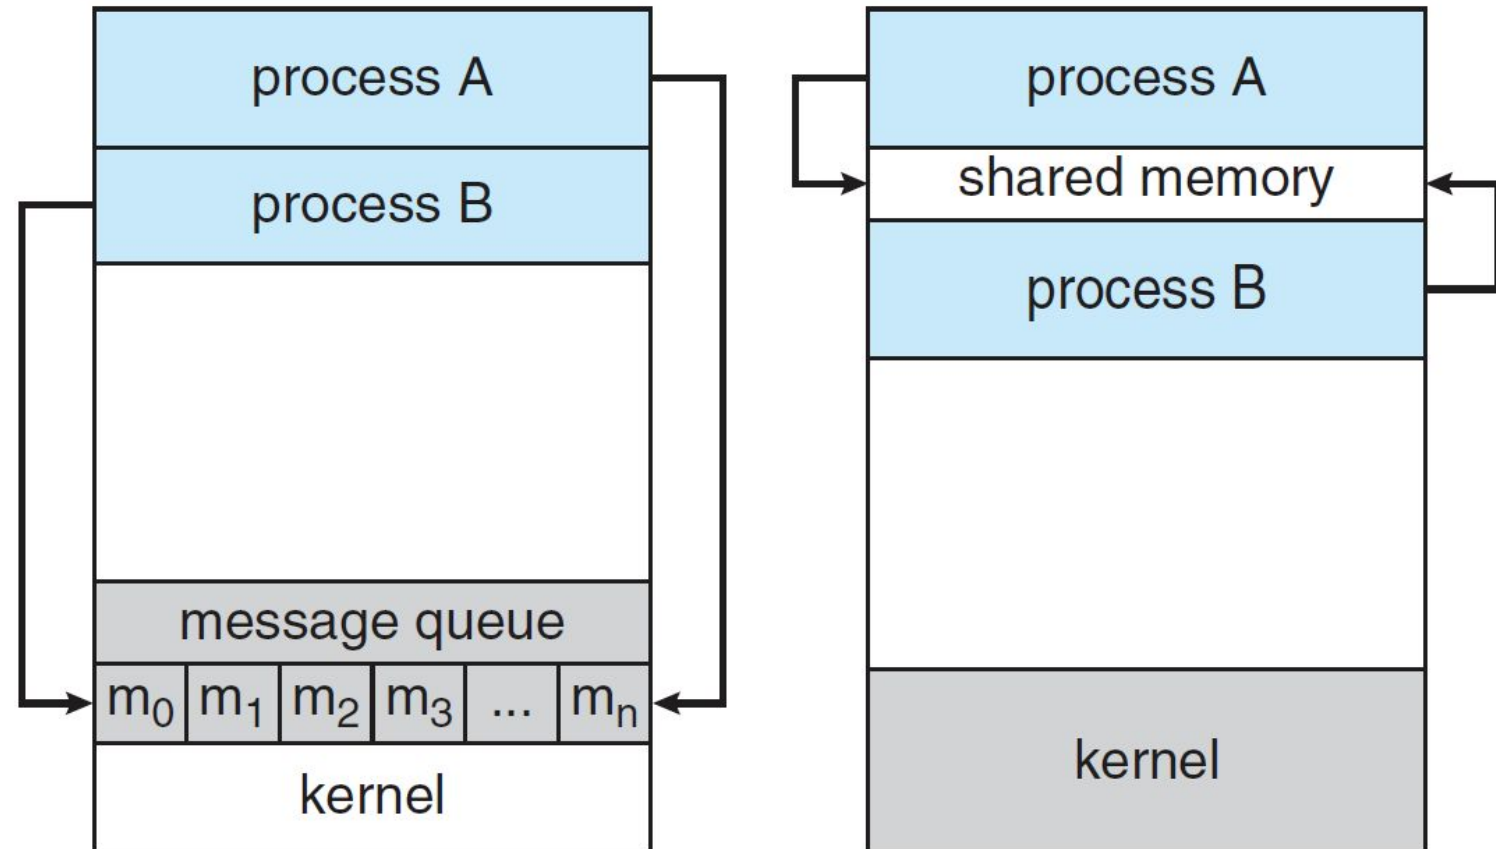
- Orphan process
 - Process whose parent has exited
 - Could litter the process table if not handled
 - On UNIX and Linux orphans are adopted by the `init` process
 - `init` periodically calls `wait()`

Interprocess communication (IPC)

- **Independent processes** – cannot affect (or be affected by) other processes
 - Do not share data with other processes
- **Cooperating processes** – opposite of independent processes
 - **Information sharing**
 - **Computation speedup** – assuming multiprocessor system
 - **Modularity**
 - **Convenience** – single user running processes concurrently
 - Require Interprocess communication (IPC)

Models of Interprocess communication

- Shared memory
- Message passing



IPC: shared memory

- Communicating processes establish a shared memory region
 - Belongs to the initiator of communication
 - OS does not arbitrate access to the shared region by the communicating processes

Classic example: the producer-consumer problem

- Producer – generates information
- Consumer – consumes the information
- Metaphor for client-server systems

Producer-consumer – Shared memory

- Shared memory contents

```
#define BUFFER_SIZE 10
typedef struct {
    . . .
} item;

item buffer[BUFFER_SIZE];
int in = 0;
int out = 0;
```

Producer

```
item next_produced;

while (true) {
    /* produce an item in next produced */
    while (((in + 1) % BUFFER_SIZE) == out)
        ; /* do nothing */

    buffer[in] = next_produced;
    in = (in + 1) % BUFFER_SIZE;
}
```

Consumer

```
item next consumed;

while (true) {
    while (in == out)
        ; /* do nothing */

    next consumed = buffer[out];
    out = (out + 1) % BUFFER SIZE;
    /* consume the item in next consumed */
}
```


Message passing

- Logical communication link
 - Direct vs. indirect communication
 - Synchronous vs. asynchronous communication
 - Automatic vs. explicit buffering
- **send(message)**
- **receive(message)**
- Communicating processes have names (IDs)

Message passing implementation questions

- How are links established?
- Can a link be associated with more than two processes?
- How many links between every pair of communicating processes?
- What is the capacity of a link?
- Is the size of a message that the link can accommodate fixed or variable?
- Is a link unidirectional or bi-directional?

Message passing – direct communication

- Sender/receiver name the other process
 - `send(P, message)`
 - `receive(Q, message)`
- 
- A diagram consisting of a rounded rectangle with the text "Symmetric addressing" inside. A large curly brace on the left side of the rectangle groups the two code snippets from the first bullet point above it: `send(P, message)` and `receive(Q, message)`.
- Asymmetric addressing – receiver does not name sender
 - `send(P, message)`
 - `receive(id, message)` – receive from any process
 - Properties of communication link
 - link established automatically – assuming the process know each other's ID
 - link associated with exactly two processes
 - Only one link between processes
 - Can be either unidirectional or bidirectional


Direct Communication – disadvantage

- Communicating processes are tightly coupled
 - process IDs are hard coded in send/receive calls
 - if process ID changes, the code of all processes that communicate with it must replace references to the old ID

Message passing – Indirect communication

- Use mailboxes (also called ports)
 - **send(A, message)**
 - **receive(A, message)**
- Link properties
 - communicating processes must share a mailbox
 - link may be associated by more than two processes
 - May be multiple links between two processes – each via a separate box
 - Unidirectional/bidirectional

Indirect communication links



Message passing – indirect communication

- Operations
 - Create/destroy mailbox
 - Send/receive messages through mailbox
- Mailbox sharing
 - *P1, P2 and P3* share mailbox *A*
 - *P1* sends; *P2* and *P3* receive
 - Who gets the message?
- Solutions
 - Allow a link to associate at most two processes
 - Allow only one process at a time to execute a receive operation
 - System arbitrarily selects the receiver. Sender is notified who the receiver was

Synchronization

- Message passing may be either blocking or non-blocking
- **Blocking** operations are **synchronous**
 - **Blocking send** – sender blocks until the message is received
 - **Blocking receive** – receiver blocks until a message is available
- **Non-blocking** operations are **asynchronous**
 - **Non-blocking send** – sender sends the message and continue
 - **Non-blocking receive** – receiver receives a valid message or null
- All combinations are possible
 - blocking send + blocking receive = rendezvous

Producer-consumer rendezvous

Producer

```
message next_produced;  
  
while (true) { /* produce an item in next produced */  
    send(next_produced);  
}
```

Consumer

```
message next_consumed;  
  
while (true) {  
    receive(next_consumed);  
    /* consume the item in next consumed */  
}
```

Buffering

- Queue of messages attached to the link
- Implemented in one of three ways
 - **Zero capacity** – 0 messages
 - Sender must wait for receiver (rendezvous)
 - **Bounded capacity** – finite length (n messages)
 - Sender must wait if link full
 - **Unbounded capacity** – infinite length
 - Sender never waits