

# Synchronization

# Motivation

- Concurrent processes accessing shared data
  - May result in data inconsistency
- Example reconsider the producer-consumer problem
  - Simplifying assumption: A single processor system
  - Remember the circular-array implementation of buffer
  - Want to use full buffer
    - Use a shared variable **counter** which keeps track of number of items in the buffer

# Producer

```
while (true) {  
    /* produce an item in next produced */  
    while (counter == BUFFER SIZE)  
        ; /* do nothing */  
  
    buffer[in] = next produced;  
    in = (in + 1) % BUFFER SIZE;  
    counter++;  
}
```

# Consumer

```
while (true) {  
    while (counter == 0)  
        ; /* do nothing */  
  
    next consumed = buffer[out];  
    out = (out + 1) % BUFFER SIZE;  
    counter--;  
    /* consume the item in next consumed */  
}
```

# Implementation of counter++ and counter--

*register1* = **counter**  
*register1* = *register1* + 1  
**counter** = *register1*



**counter++**

*register2* = **counter**  
*register2* = *register2* - 1  
**counter** = *register2*



**counter--**

# Race conditions

- Assuming the current value of counter is 5
  - What will be the value if both producer and consumer run concurrently?
  - Answer: 4, 5 or 6
  - Why?

# Race conditions

*T0: producer* execute **register1 = counter**      {*register1* = 5}  
*T1: producer* execute **register1 = register1 + 1** {*register1* = 6}  
*T2: consumer* execute **register2 = counter**      {*register2* = 5}  
*T3: consumer* execute **register2 = register2 - 1** {*register2* = 4}  
*T4: producer* execute **counter = register1**      {*counter* = 6}  
*T5: consumer* execute **counter = register2**      {*counter* = 4}

# The Critical Section Problem

- Each process has a so called critical section
  - Changes the value of shared data
- Critical section problem: design a protocol that processes can use to cooperate



# Structure of a process

do {

*entry section*

critical section

*exit section*

remainder section

} while (true);

# Constraints on solutions to the Critical Section Problem

- Mutual exclusion
  - Exactly one process may be executing in its critical section
- Progress
  - Only processing requesting entry into critical section may be considered for entry
- Bounded waiting
  - Limit on the number of times entry request may be deferred

# Race conditions in the kernel

- Preemptive kernels
  - Possibility of race conditions – need an avoidance protocol
  - Better system performance – a kernel process is not allowed to hold on to the CPU for too long
- Nonpreemptive kernels
  - No possibility of race conditions on single processor systems
  - Poor system performance

# Peterson's Solution

- Restricted to two process that take turns in executing their critical sections

```
do {  
    flag[i] = true;  
    turn = j;  
    while (flag[j] && turn == j);  
    critical section  
    flag[i] = false;  
    remainder section  
} while (true);
```

# Synchronization hardware

- Locking
- One possibility is to disable interrupts
  - Only works for single processor systems
  - Inefficient for multicore systems – message has to be passed to all processors
  - Aversely affects system clock, if it's update is interrupt based
- Modern CPU architectures provide two atomic (uninterruptible) operations
  - **test\_and\_set()**
  - **compare\_and\_swap()**
  - Can use these to solve the critical section problem

# test\_and\_set()

```
boolean test_and_set(boolean *target) {  
    boolean rv = *target;  
    *target = true;  
    return rv;  
}
```

- target is a pointer (reference) parameter
  - Changes to target are reflected in the actual parameter

# test\_and\_set()

do {

```
while (test_and_set(&lock))  
    ; /* do nothing */
```

```
    /* critical section */
```

```
    lock = false;
```

```
    /* remainder section */
```

```
} while (true);
```

# compare\_and\_swap()

```
int compare_and_swap(int *value, int expected, int new_value) {  
    int temp = *value;  
  
    if (*value == expected)  
        *value = new_value;  
  
    return temp;  
}
```



# compare\_and\_swap()

do {

```
while (compare_and_swap(&lock, 0, 1) != 0)  
    ; /* do nothing */
```

```
    /* critical section */
```

```
    lock = 0;
```

```
    /* remainder section */
```

```
} while (true);
```

# Synchronization hardware

- The two algorithms just presented
  - Guarantee mutual exclusion
  - Do not guarantee bounded waiting
- Next algorithm guarantees both

```
do {  
    waiting[i] = true;  
    key = true;  
    while (waiting[i] && key)  
        key = test_and_set(&lock);  
    waiting[i] = false;  
    /* critical section */  
    j = (i + 1) % n;  
    while ((j != i) && !waiting[j])  
        j = (j + 1) % n;  
    if (j == i)  
        lock = false;  
    else  
        waiting[j] = false;  
    /* remainder section */  
} while (true);
```

# Mutex locks

- **Mutex** (= **mutual exclusion**) lock
  - Software mechanism available to application programmers
  - **acquire()** – process must successfully execute this before entering CS
  - **release()** – process must execute this on exit to free the lock
  - **available** – Boolean indicating whether the lock is available

# Mutex locks

```
acquire() {  
    while (!available)  
        ; /* busy wait */  
  
    available = false;;  
}
```

```
release() {  
    available = true;  
}
```

# Mutex locks

do {

*acquire lock*

critical section

*release lock*

remainder section

} while (true);

# Mutex locks

- Disadvantage
  - busy waiting – spinlock
    - Wastes CPU cycles – a problem in a multiprogramming environment
- Spinlocks
  - Can be useful for short waiting time
    - When context-switching would consume more time

# Semaphores

- Semaphore is an integer + two atomic operations
  - wait()
  - signal()
- Binary semaphore – domain is  $\{0, 1\}$ 
  - Same functionality as a mutex lock
- Counting semaphore – unrestricted domain
  - Controls to a resource with finite instances



# Semaphores

```
wait(S) {  
    while (S <= 0)  
        ; // busy wait
```

```
    S--;  
}
```

```
signal(S) {  
    S++;  
}
```

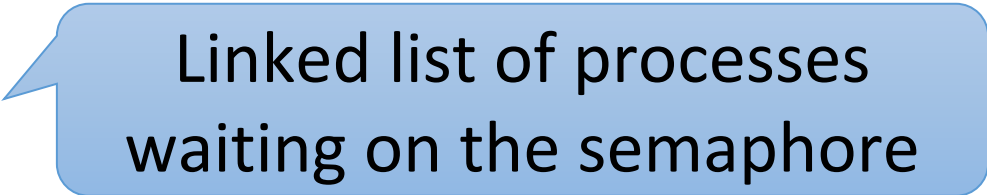
# Binary semaphore – use case

- Synchronizing two processes
  - Say P2 must statement S2 only after P1 executes statement S1
  - Initialize semaphore synch to 0

```
P1{  
    ...  
    S1;  
    signal(synch)  
    ...  
p2{  
    ...  
    wait(synch);  
    S2
```

# Semaphore + blocking **wait()**

```
typedef struct {  
    int value;  
    struct process *list;  
} semaphore;
```



Linked list of processes  
waiting on the semaphore

# Semaphore + blocking **wait()**

```
wait(semaphore *S) {  
    S->value--;  
  
    if (S->value < 0) {  
        add this process to S->list;  
        block();  
    }  
}
```

# Semaphore + blocking **wait()**

```
signal(semaphore *S) {  
    S->value++;  
  
    if (S->value <= 0) {  
        remove a process P from S->list;  
        wakeup(P);  
    }  
}
```

# Semaphores – blocking **wait()**

- Both **wait()** and **signal()** must be executed atomically
  - Critical sections!
- Atomicity can be guaranteed by
  - Disabling interrupts – sensible only on single processor systems
  - Use **test\_and\_set()** or **compare\_and\_swap()**
- Blocking wait does not completely eliminate busy waiting

# Deadlocks

- **Defn:** A set of process (two or more processes in the set) are deadlocked if each process in the set is waiting for another process in the set to execute `signal()`
  - Since the process that's supposed to execute `signal` is blocked, all processes in the set will wait indefinitely.

# Deadlock example

$P_0$

`wait(S);`

`wait(Q);`

•

•

•

`signal(S);`

`signal(Q);`

$P_1$

`wait(Q);`

`wait(S);`

•

•

•

`signal(Q);`

`signal(S);`



# Indefinite blocking (starvation)

- A process may be indefinitely blocked is the queue on a semaphore is processed in, say, LIFO

# Priority Inversion

- Occurs on systems implementing more than two levels of priority
- Example: Assume process L, M and H have low, medium and high priority resp.
  - L is currently executing and using resource R
  - H becomes runnable and also requests R – it blocks since R is locked, so L continues
  - M becomes runnable and preempts L
  - M has been dispatched while H (with a higher priority) has to wait for a resource being held by L

# Priority Inversion - prevention

- Restrict priority levels to 2
  - Not practical on most systems
- Priority inheritance
  - Lower priority process holding a resource required by higher priority process temporarily adopts the higher priority

# Classical synchronization problems

- Bounded buffer
- Readers-writers
- Dining philosophers

# Bounded buffer problem

- Producer-consumer problem using a bounded buffer
- Shared data

```
int n;  
semaphore mutex = 1;  
semaphore empty = n;  
semaphore full = 0
```

# Producer

```
do {  
    . . .  
    /* produce an item in next produced */  
    . . .  
    wait(empty);  
    wait(mutex);  
    . . .  
    /* add next produced to the buffer */  
    . . .  
    signal(mutex);  
    signal(full);  
} while (true);
```

# Consumer

```
do {  
    wait(full);  
    wait(mutex);  
    . . .  
    /* remove an item from buffer to next consumed */  
    . . .  
    signal(mutex);  
    signal(empty);  
    . . .  
    /* consume the item in next consumed */  
    . . .  
} while (true);
```

# Readers-writers problem

- First readers-writers problem
  - One writer at a time – no readers permitted
  - Multiple readers at a time
  - Readers must not wait for writer to start writing
- Second readers-writers
  - Writer must write as soon as possible – no new reader admitted as soon as writer is ready to write
- Applications:
  - Can distinguish between readers and writers
  - More readers than writers



# First readers-writers problem – shared data

```
semaphore rw_mutex = 1;  
semaphore mutex = 1;  
int read count = 0;
```

# First readers-writers problem - writer


```
do {  
    wait(rw_mutex);  
    . . .  
    /* writing is performed */  
    . . .  
    signal(rw_mutex);  
} while (true);
```

# First readers-writers problem - reader

```
do {
    wait(mutex);
    read count++;
    if (read count == 1)
        wait(rw_mutex);
    signal(mutex);

    . . .
    /* reading is performed */
    . . .
    wait(mutex);
    read count--;
    if (read count == 0)
        signal(rw_mutex);
    signal(mutex);
} while (true);
```

# Dining philosophers problem



# Dining philosophers problem

```
do {  
    wait(chopstick[i]);  
    wait(chopstick[(i+1) % 5]);  
    . . .  
    /* eat for awhile */  
    . . .  
    signal(chopstick[i]);  
    signal(chopstick[(i+1) % 5]);  
    . . .  
    /* think for awhile */  
    . . .  
} while (true);
```

# How to deal with deadlocks

- Allows only 4 philosophers to attempt to pick up a chopstick
- Allow a philosopher to attempt to pick up chopsticks only if both are free – critical section
- Asymmetric solution

# Weakness of semaphores

- Non-adherence to the **wait()-signal()** sequence by processes
  - Deliberate or accidental

```
signal(mutex);  
...  
critical section  
...  
wait(mutex);
```

- Broken mutual exclusion

# Weakness of semaphores

- Non-adherence to the **wait()-signal()** sequence by processes
  - Deliberate or accidental

```
wait(mutex);
```

```
...
```

```
critical section
```

```
...
```

```
wait(mutex);
```


- Deadlock



# Monitors

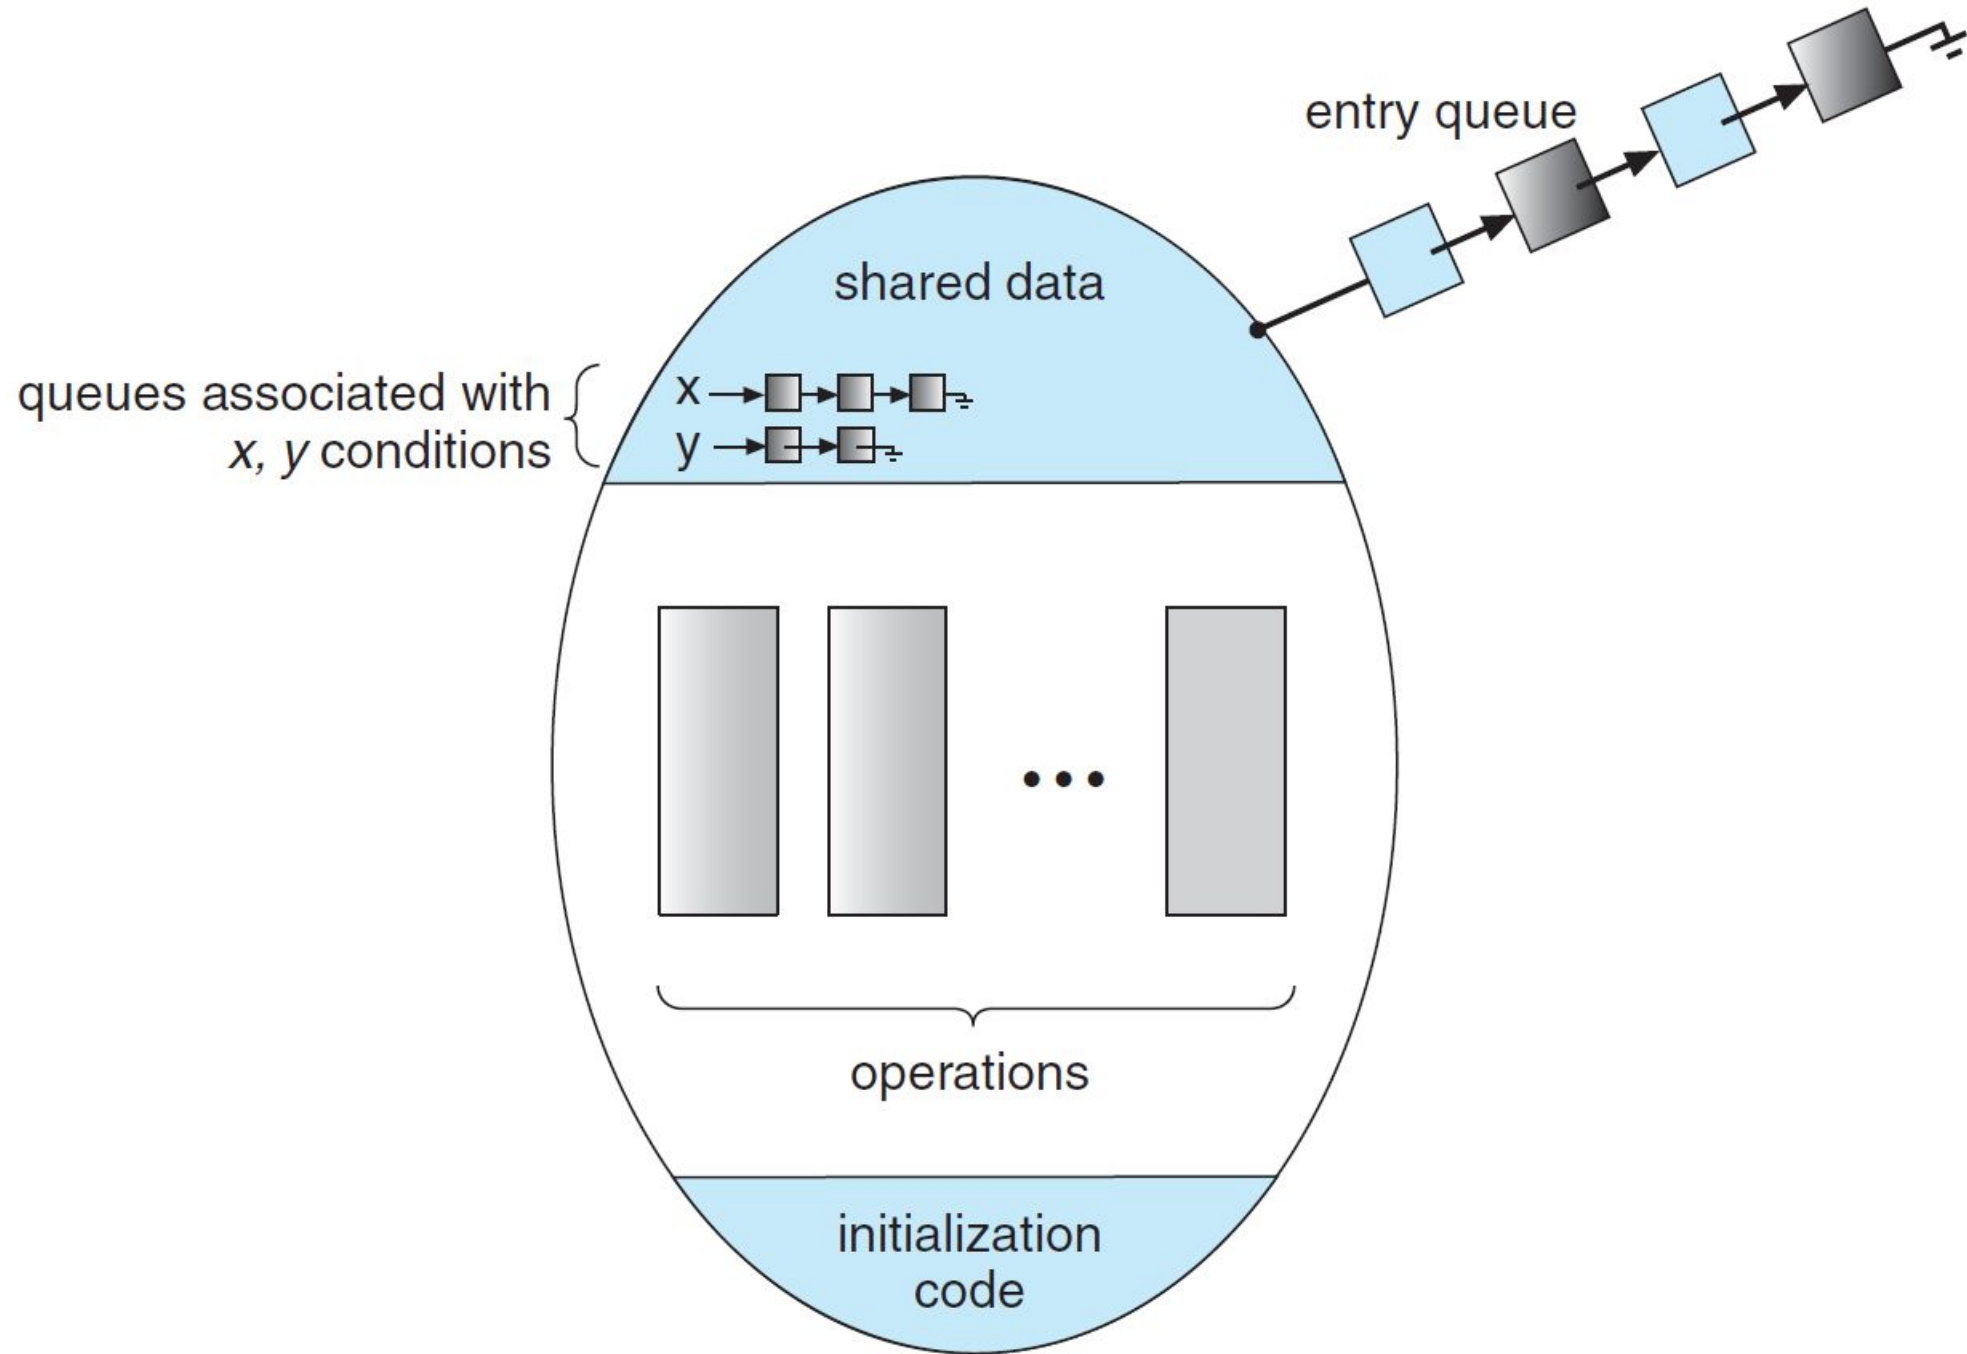
- Programming language support for synchronization
- Solves the semaphore problems
- Available in languages like Java, C#
- Monitor ADT

```
monitor monitor_name{  
    /* shared variable declarations */  
  
    function P1 ( . . . ) {  
        . . .  
    }  
  
    function P2 ( . . . ) {  
        . . .  
    }  
    .  
    .  
    .  
    function Pn ( . . . ) {  
        . . .  
    }  
  
    initialization_code ( . . . ) {  
        . . .  
    }  
}
```



# Monitors

- **condition** type
  - supports two operations
    - **x.wait()**
    - **x.signal()**
- What happens when process P signals process Q?
  - Either **signal and wait** – P must wait for Q to exit monitor or wait
  - Or **signal and continue** – Q must still wait for P to exit monitor or wait



# Dining philosophers solution using monitors

```
monitor DiningPhilosophers{
    enum {THINKING, HUNGRY, EATING} state[5];
    condition self[5];

    void pickup(int i) {
        state[i] = HUNGRY;
        test(i);
        if (state[i] != EATING)
            self[i].wait();
    }

    void putdown(int i) {
        state[i] = THINKING;
        test((i + 4) % 5);
        test((i + 1) % 5);
    }
    ...
}
```

# Dining philosophers solution using monitors

```
monitor DiningPhilosophers{
    ...
    void test(int i) {
        if ((state[(i + 4) % 5] != EATING) &&
            (state[i] == HUNGRY) &&
            (state[(i + 1) % 5] != EATING)) {
            state[i] = EATING;
            self[i].signal();
        }
    }

    initialization_code() {
        for (int i = 0; i < 5; i++)
            state[i] = THINKING;
    }
}
```

# Dining philosophers solution using monitors

- No two neighbors eat at the same time
- No deadlocks
- Starvation may still occur